

DSA Top-K Indexer: Process Write-Up

Author: Yury Kirpichev / Team Wombat Track: `dsa_topk_indexer_fp8_h64_d128_topk2048_ps64` Submission tags: `submission-v1` through `submission-v11` Final solution: `dsa_topk_indexer_fp8_b200_v3`

Headline Summary

Field	Value
Track	<code>dsa_topk_indexer_fp8_h64_d128_topk2048_ps64</code>
Hardware	NVIDIA B200 / <code>sm_100a</code> (FP8 E4M3 x FP8 E4M3 -> FP32, FP16 logits)
Tag	<code>submission-v11</code> (commit <code>31b8f71</code>)
Workloads	128 / 128 PASSED
vs naive ref	964x mean speedup (<i>unconfirmed — see caveat below</i>)
vs FlashInfer/DG	38.4x mean (8.3x worst, 72.3x best) (<i>unconfirmed — see caveat below</i>)
Aggregate latency	mean 7.83 us, p50 2.40 us, p95 17.80 us (<i>unconfirmed — see caveat below</i>)
Source files	<code>solution/python/solution.py</code> , <code>solution/python/kernel.cu</code> , <code>solution/python/tcgen05_ptx.h</code> , <code>solution/python/umma_desc.h</code>
Methodology	Agent-assisted (Cursor + LLM agents); all decisions and commits owned by the human author. See "AI / Agent-Assisted Development Disclosure" after the References.


Caveat — comparison numbers are unconfirmed. The headline speedups vs the FlashInfer / DeepGEMM baseline (38.4x mean) and vs the naive Python reference (964x mean) are taken from `reports/submission-v10.md` and the raw log in `reports/submission-v10-raw-cupti.log`. It is **unclear at the time of writing whether the comparison numbers were end-to-end** `cupti-python-timed` **using the same harness configuration the official evaluator uses**, or whether some of the relative numbers were derived from a mix of `cupti` and `CUDA-event` measurements.

Problem, Constraints, and Development Method


The DSA top-K indexer computes, for each batch row `b`:

- FP8 dot products:** `S[h] = sum_d Q[b,h,d] * K[page,slot,d]` over 64 heads and 128 dimensions, with FP8 E4M3 inputs and FP32 accumulation.
- Weighted ReLU logits:** `logit[b,t] = scale[page,slot] * sum_h ReLU(S[h]) * weights[b,h]` for every context position `t` in `[0, seq_lens[b))`.
- Top-K selection:** the 2048 highest-scoring positions.
- Page-table transform:** converting token-level indices to physical page offsets via `block_table`.

The KV cache is packed in FP8 with per-slot FP32 scales at a stride of 132 bytes per slot (128 FP8 values + 4 bytes FP32 scale). Pages are 64 slots each. The operator receives pre-allocated output tensors in destination-passing style.



The practical challenge was that the 128 contest workloads span a wide range of shapes: batch sizes 1-16, `max_num_pages` from 5 to 95+, meaning some workloads have only ~320 tokens while others exceed 6000. For the small workloads, host launch overhead and TMEM allocation cost dominate; for the large ones, HBM bandwidth for K-cache fetches and the TMEM-to-register readout path are the bottleneck. A single kernel cannot serve both regimes well.



The development followed an experiment-driven loop. Each optimization was benchmarked on Modal B200 with `cupiti-python` timing, and either kept or reverted based on measured numbers. Negative results were tracked alongside positive ones, which proved essential for understanding the actual critical path.

Challenges: Evaluation Framework (Before and After PR #354)

The single largest constraint on the development timeline was the evaluation framework itself. Before [flashinfer-bench PR #354](#) (merged April 10, 2026), there was no dedicated `DsaTopKIndexerEvaluator`. The default evaluator compared **raw output index vectors** element-wise, which had two consequences that shaped the first six weeks of work:

1. **Tie-breaking forced determinism.** When multiple context positions share the same logit value, any subset of the tied positions is a correct top-K answer. But raw index comparison rejects every ordering that differs from the reference, even if the selected values are identical. This meant that any custom top-K implementation — radix select, CUB sort with nondeterministic key ordering, or any `atomicAdd`-based emit — would produce valid top-K sets that the evaluator rejected as `INCORRECT_NUMERICAL`. In

practice, **the only top-K path that reliably passed the old evaluator was `torch::topk`**, which is deterministic and matches the reference's tie-breaking order exactly.

2. **FP8 packing mismatch.** The default evaluator's random `k_index_cache_fp8` inputs were not packed in the deepgemm FP8 format that the FlashInfer baseline expected, causing false negatives where correct FlashInfer kernels were themselves rejected.

What this meant for development (submissions v1-v4, March 29 - April 6).

The early submissions were constrained to ATen-level operations that preserved bit-exact index agreement with the reference:

- `submission-v1` (Mar 29): fused FP8 page-gather + dequant CUDA kernel, then `torch::bmm` for Q·K, ATen `relu` + weighted sum, and `torch::topk` for selection. ~5.6x vs naive reference, 128/128 PASSED.
- `submission-v2` (Mar 30): batched page transform, minor ATen reordering. ~6.8x, 128/128 PASSED.
- `submission-v3` (Mar 31): in-place `relu` + `mul` on logits. ~6.6x.
- `submission-v4` (Apr 6): buffer reuse for `sum_out` and `topk_vals`. ~7.4x.

Several more aggressive optimizations were attempted and reverted during this period because they broke bit-exact index agreement: fused Q·K dot products, per-row score building without the full `[B,H,S]` weighted tensor, BF16 reduction paths, and chunked GEMM with running top-K. Each was numerically correct (produced a valid top-K set) but failed the raw-index evaluator.

After PR #354 (April 10). The new `DsaTopkIndexerEvaluator` compares **sorted value vectors** instead of raw indices, correctly handling tie-breaking and rounding differences. It also includes vectorized index validation (duplicates, out-of-range, block-table reachability) and a custom `build_baseline` with correct FP8 packing.

This immediately unblocked the entire custom-kernel path. Within ten days of the fix (April 11-24), the solution was rewritten from scratch:


- April 11-12: pure PyTorch FP8 dequant + bmm indexer (no deepgemm dependency), then chunked GEMM + running top-K.
- April 19-20: custom CUDA top-K with CUB radix sort (`abserr=0`, `relerr=0`); then FP8 MMA logits via `mma.sync.m16n8k32.e4m3 PTX` + persistent-queue scheduling. Mean 12x vs naive, 128/128.
- April 20-21: full SM100a UMMA rewrite with `tcgen05.mma.cta_group:1`, radix-select top-K, size-aware dispatch (`submission-v7`, `v8`). ~16 us mean, 3.5x vs FlashInfer.
- April 21-24: persistent CTAs, warp specialization, Stage-2 SMEM cache, dispatch re-tuning, CUDA graph cache (`submission-v9` through `v11`). Final: 7.83 us mean, 38.4x vs FlashInfer.

In summary, the pre-#354 evaluator constraint kept the solution at the ATen/ `torch::topk` level for the first six weeks (~7x speedup ceiling), while the post-#354 evaluator enabled the full custom CUDA pipeline that reached the final 38.4x in the remaining two weeks.

Speedup trajectory across submissions

The chart below illustrates the speedup ceiling jump immediately after PR #354. Numbers are rounded; pre-#354 submissions are speedup vs naive Python reference on the early benchmark suite (preliminary), post-#354 are vs the FlashInfer baseline reported by `flashinfer-bench` (also preliminary; see caveat above).


Mean speedup vs naive reference, by submission



The flat plateau at v1-v4 (5-7x) corresponds to the period when the old evaluator rejected anything beyond `ATen + torch::topk`. The sharp rise at v7-v8 follows PR #354 and the full custom-CUDA + UMMA + radix-select rewrite. The final ~13% from v9 to v11 comes from CUDA-graph host-overhead elimination on the small-workload bucket.

Chronological Optimization Story

The repo history clusters into six phases, split by the PR #354 watershed:



Phase 0 (`submission-v1` through `submission-v4` , March 29 - April 6) operated under the old evaluator's raw index-comparison constraint. The kernel used a custom CUDA FP8 page-gather + dequant kernel for Stage 1, then `torch::bmm` for Q·K dot products, ATen `relu` + weighted sum for logit computation, and `torch::topk` for selection — the only top-K implementation that produced bit-exact index agreement with the reference. Multiple attempts at more aggressive optimizations (fused Q·K, per-row scoring, BF16 reduction, chunked GEMM + running top-K) were reverted because they broke index-level agreement even when the selected values were correct. The ceiling under this constraint was ~7.4x vs the naive reference.

Phase 1 (`submission-v7` , `submission-v8` , April 19-21) became possible after PR #354 removed the bit-exact index constraint. The solution was rewritten from scratch as a fully custom CUDA kernel. The key design was to

use Blackwell's SM100a UMMA (`tcgen05.mma.cta_group::1.kind::f8f6f4`) for the FP8 dot products, mapping the `[64 slots, 64 heads]` tile onto UMMA's `M=128, N=64, K=32` shape with 4 K-iterations. Each CTA handled one page (64 KV slots), loaded Q and K into shared memory with 128B swizzled layout, allocated TMEM, issued the UMMA, read back the FP32 accumulator, computed ReLU-weighted logits, and stored FP16 results. Stage 2 used a 2-round radix select for top-K — an approach that would have been impossible under the old evaluator. By `submission-v8`, the kernel achieved ~16 us mean across all 128 workloads, a ~5x jump over the Phase 0 ceiling.

Phase 2 applied structural optimizations. Persistent CTAs amortized Q loading and TMEM allocation across tile-pairs. A 3-stage `cp.async K` pipeline overlapped HBM fetches with UMMA compute. Warp specialization split 256 threads into a producer warpgroup (128 threads, 40 registers, handles `cp.async K` prefetch) and a math warpgroup (128 threads, 232 registers, handles UMMA issue, TMEM readout, ReLU-weighted sum, and emit). Per-stage `K_ready / K_done` mbarriers drive the producer-math handoff. A critical correctness fix discovered during bring-up: `cp.async.wait_group` is per-thread, so a producer-warpgroup `bar.sync` was needed before the single-thread `mbarrier.arrive` on `K_ready[buf]` to ensure all threads' `cp.async` writes are visible.

Several ideas in Phase 2 were tried and rejected based on measured regressions:

Lever	Outcome	Why
TMEM double-buffer + MMA-issue overlap	+3.6-5.2%	MMA compute is not on the critical path; <code>tcgen05_wait_ld</code> dominates
kKVStages 3 -> 4	Flat/+0.9%	Producer not the bottleneck
4-way accumulator split in FMA chain	Within noise	FMA chain already hidden behind <code>tcgen05_wait_ld</code> latency
Block-scaled UMMA (<code>mxsf8f6f4</code>)	Infeasible	Data format uses per-row FP32 scales, not per-K-block UE8M0
Host-side <code>seq_lens.max()</code> sync	+20 us	<code>cudaStreamSynchronize</code> costs ~15-20 us on B200, more than any savings


The combined evidence identified the per-tile bottleneck as the TMEM-to-register transfer (`tcgen05_ld_32x32b_x64_b32 + tcgen05_wait_ld`), which is tied to `kUMMA_N = 64` columns. The only remaining structural lever is 2-CTA cluster UMMA, which was deferred due to infrastructure complexity.

Phase 3 ported micro-optimizations from a parallel branch. The Stage-2 ordered-u16 SMEM cache converts logits to comparison-ordered uint16 once on round 0 and reads from shared memory on subsequent passes, saving 4x HBM reads. The fast-path kernel was vectorized to 128 threads with int4 stores. The persistent-kernel dispatch threshold was re-tuned from `max_num_pages >= 40` to `>= 64`, recovering 8.9% on the 40-63 page bucket where the short kernel's grid-per-tile layout saturates B200's 148 SMs faster than the persistent pipeline.

Phase 4 (`submission-v10`, `submission-v11`) added a hand-rolled CUDA graph cache. The benchmark harness `do_bench` clones input tensors on every iteration, so pointers change but shapes (and the entire dispatch plan) stay fixed. For small workloads where the kernel itself runs in 2-10 us, the ~2.5 us per `cudaLaunchKernel` dominated. The graph cache keys by shape only (stream, dispatch path, grid-defining scalars), captures via `cudaStreamBeginCapture` on a dedicated thread-local stream, and on cache hits uses `cudaGraphExecUpdate` to splice fresh pointers into the cached exec in place. This collapsed the fast-path bucket from 6.5 us to 2.3 us mean.

Final Kernel Design

The final kernel consists of a host dispatcher, three Stage-1 kernel variants, a Stage-2 radix top-K kernel, and a fast-path bypass, all wrapped in a CUDA graph cache.




Fast path ($\text{max_num_pages} \leq 32$). When the entire paged context fits within $K=2048$, every position is in the top-K. The output is the block-table-transformed indices $[0, \text{seq_len})$ padded with -1, independent of Q, K, and weights. A 128-thread kernel with vectorized int4 stores handles this in ~ 2.3 us including graph replay overhead.

Short kernel ($33 \leq \text{max_num_pages} < 64$). One CTA per page, 128 threads. Loads Q and K into swizzled SMEM, allocates TMEM, issues UMMA with 4 K-iterations, reads back FP32 accumulator, computes ReLU-weighted logits, emits FP16. Simple grid layout with minimal setup cost.

Warp-specialized persistent kernel ($\text{max_num_pages} \geq 64$). 256 threads split into producer (warps 4-7, 40 regs) and math (warps 0-3, 232 regs) warpgroups. Producer handles 3-stage `cp.async K` pipeline; math handles UMMA, TMEM readout, and emit. Q loaded once, TMEM allocated once per CTA. Grid is $(\text{num_splits}, B)$ with `tiles_per_cta` tile-pairs per CTA.

The producer/math handoff is mediated by per-stage `K_ready[buf]` and `K_done[buf]` mbarriers, enabling a 3-deep `cp.async` pipeline that overlaps HBM K-fetches with UMMA compute and TMEM readout:



Stage-2 radix top-K. Two-round radix select over 16-bit ordered keys, finding the pivot value in $O(\text{seq_len})$ work. When `seq_len` ≤ 16384 (all workloads in this benchmark), ordered keys are cached in a 32 KB SMEM array on round 0 and re-read from cache on subsequent passes. Elements greater than the pivot are emitted first, then elements equal to the pivot fill the remaining slots. A final pass applies the block-table transform.

CUDA graph cache. A 32-slot LRU keyed by `(stream, dispatch_path, B, max_num_pages, tiles_per_cta, num_splits, max_len)`. Each call captures a fresh graph with current pointers and tries `cudaGraphExecUpdate` against the cached exec. On cache miss, `cudaGraphInstantiate` creates a new exec. Falls back to direct launch if `cudaStreamBeginCapture` fails.

Negative results. The most instructive finding was that the per-tile critical path in the persistent kernel is dominated by the TMEM-to-register transfer (`tcgen05_wait_ld`), not by MMA compute, producer prefetch latency, or the FMA accumulation chain. Every attempted micro-optimization that added work to the post-`wait_ld` emit path (ordered-u16 conversion in Stage 1, fused block-table transform in emit, skipped pad writes) regressed the 40-63 page bucket by 0.6-3.7% because it lengthened the critical path between successive TMEM loads. The only structural way to reduce `wait_ld` cost is to shrink `KUMMA_N` via 2-CTA cluster UMMA, which was deferred due to infrastructure complexity (new PTX wrappers, cluster launch, DSMEM, cluster-scoped mbarriers).

Validation, Results, and Reproducibility

Correctness was validated continuously via the benchmark harness. The final submission passes all 128 workloads with the `DsaTopkIndexerEvaluator` (post-PR #354).

Tools and Languages

Tool / Language	Role
CUDA C++ (sm_100a)	All kernels: UMMA via inline PTX, cp.async pipelines, radix top-K
<code>torch.utils.cpp_extension.load()</code>	JIT compilation with <code>-gencode arch=compute_100a,code=sm_100a</code>
Inline PTX (<code>tcgen05_ptx.h</code>)	TMEM alloc/dealloc, UMMA issue/commit/wait, cp.async, mbarrier, warpgroup reg reconfig
CUTLASS bitfield layouts (<code>umma_desc.h</code>)	SmemDescriptor and InstrDescriptor for UMMA, vendored for self-containment
Python (<code>solution.py</code>)	Thin wrapper: JIT-compiles and calls the extension
Modal	Cloud B200 access for benchmarking
<code>cupiti-python</code>	GPU-side timing (matches official evaluation methodology)
Cursor + LLM agents	Agent-assisted development (see disclosure below)

References

Submitted source

1. Final kernel implementation - `[solution/python/kernel.cu](solution/python/kernel.cu)` .
2. Python entry point and JIT compile flags - `[solution/python/solution.py](solution/python/solution.py)` .
3. Inline PTX wrappers for SM100a - `[solution/python/tcgen05_ptx.h](solution/python/tcgen05_ptx.h)` .
4. UMMA descriptor layouts (vendored from CUTLASS) - `[solution/python/umma_desc.h](solution/python/umma_desc.h)` .
5. Submission manifest - `[config.toml](config.toml)` .

Contest documentation

1. Contest rules - `[README.md](README.md)` .
2. FAQ (allowed optimizations, `torch.utils.cpp_extension` , custom compile flags) - `[FAQ.md](FAQ.md)` .
3. Evaluation environment - `[EVALUATION.md](EVALUATION.md)` .

Performance artifacts

1. Final cupiti-timed comparison report - `[reports/submission-v10.md](reports/submission-v10.md)` .
2. Raw cupiti-timed log - `[reports/submission-v10-raw-cupiti.log](reports/submission-v10-raw-cupiti.log)` .
3. Optimization history and ablation results - `[IMPROVEMENTS.md](IMPROVEMENTS.md)` .
4. Future ideas and deferred optimizations - `[experiments/FUTURE_IDEAS.md](experiments/FUTURE_IDEAS.md)` .

External references

1. Flashinfer-bench PR #354: `DsaTopkIndexerEvaluator` for correct tie-breaking comparison - github.com/flashinfer-ai/flashinfer-bench/pull/354.
2. FlashAttention-4 (Dao et al., arXiv 2603.05451, Mar 2026) - asymmetric scaling, 2-CTA MMA, TMEM intermediates, warp specialization.
3. NVIDIA CUTLASS Blackwell docs - `tcgen05.mma.kind::f8f6f4`, `cta_group=2`, `SmemDescriptor/InstrDescriptor` layouts.
4. Colfax Research - "Writing GEMM Kernels Using Tensor Memory For Blackwell GPUs."
5. DeepGEMM `sm100_fp8_paged_mqa_logits.cuh` - reference production kernel with warp specialization, TMA pipelines, persistent scheduling.
6. NVIDIA CUDA C++ Programming Guide, CUDA graphs and `cudaGraphExecUpdate` - docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#cuda-graphs.
7. PyTorch JIT C++/CUDA extensions, `torch.utils.cpp_extension.load` - pytorch.org/docs/stable/cpp_extension.html.
8. FlashInfer attention library (baseline used by the contest harness) - github.com/flashinfer-ai/flashinfer.

AI / Agent-Assisted Development Disclosure

Agent-assisted (Cursor + Claude/GPT-family LLMs); direction and acceptance criteria were human, all changes were accepted on measured B200 numbers. The agent was used as a pair-programming and refactoring tool for code generation, inline PTX wrapper authoring, and experiment scaffolding. All design decisions (kernel architecture, dispatch thresholds, pipeline depth, optimization accept/reject) were made by the human author based on benchmark measurements. The agent never had access to contest workload data or evaluation results beyond what was visible in the development logs.